

Long Quan Ton

Ho Chi Minh City, Vietnam | tonlongquanvn@gmail.com | [linkedin/in/benedictquanton](https://www.linkedin.com/in/benedictquanton) | [github/BenedictQuanTon](https://github.com/BenedictQuanTon)

EDUCATION & CERTIFICATIONS

University of Technology Sydney & HCMUT

Bachelor of Artificial Intelligence; GPA: **6.75/7.0** (Dean's List 2026)

Ho Chi Minh City, Vietnam

Expected 2028

Certifications:

- IELTS: 7.0
- Generative AI with Large Language Models
- Machine Learning Specialization

British Council
AWS/DeepLearning.AI
Stanford Online/DeepLearning.AI

EXPERIENCE

FlyRank AI

AI Engineering Intern

Remote

Aug 2026 – Present

- Architected high-throughput RESTful API endpoints and backend microservices using FastAPI, Redis, and Python, improving system response latency by **~40%** across production workflows.
- Engineered autonomous multi-agent reasoning flows using LangGraph, LangChain, and local LLMs via Ollama (Qwen 2.5 & Llama 3), accelerating task execution speeds and complex query resolution by **~65%**.
- Optimized development velocity and API spec design throughput by **~3x** through intelligent integration of Claude Opus 5 and Cursor IDE workflows into daily engineering pipelines.

Ho Chi Minh City University of Technology (HCMUT)

Undergraduate AI Research Assistant – AITechLab (ML4U)

Ho Chi Minh City

Aug 2026 – Present

- Researching **Automatic Speech Recognition (ASR)** in the Speech Recognition Team under Dr. Nguyen Duc Dung, analyzing model internals and benchmarking STT/TTS/LALMs systems for low-latency, real-time conversational voice agents, contributing to upcoming scientific publications and technical reports.
- Exploring **Small Language Models (SLMs)** for on-device inference, evaluating resource-efficiency trade-offs to enable seamless end-to-end speech-to-speech pipelines.

COMPETITIONS & PROJECTS

Weatherise – Top 10 Finalists Vietnam AI Open Hackathon

Project Lead & AI Developer

NVIDIA/Viettel/Sovico

Jun 2, 2026 – Jun 12, 2026

- Led a 4-person team to a **Top 10 finish** with a multi-agent weather decision system for Tourism, Construction & Agriculture, matching domain experts on **91%** of go/no-go calls (**120 scenarios**) at **4.8s median latency (1,247 runs)** on **8x NVIDIA H200**, using LangGraph, NeMo Agent Toolkit, and NVIDIA NIM.
- Cut forecast error by **12% MAE** vs. the best single source (**22%** vs. raw GFS, **180 days**) by fusing **7 live weather APIs** through a **10-tool MCP Server** with bias correction, outlier rejection, and an LLM arbiter, using FastAPI and Redis 7.
- Achieved **0 unsafe go-calls** across **212 red-team prompts** (94% blocked at input, the rest vetoed by a deterministic rule engine) and **87.2% context recovery** via a **4-collection RAG layer** (**~2,000 records**) behind NeMo Guardrails, using nv-embedqa-e5-v5, Qdrant, and PostgreSQL 16.

The Lantern – AssemblyAI Voice Agent Hackathon

Project Lead & AI Engineer

AssemblyAI

Sep 8, 2026 – Sep 30, 2026

- Led a 5-person team to build a real-time, multilingual AI waiter that runs a full table's service (recommendations, orders, sold-out swaps, allergy checks, live kitchen sync) at **<\$1 per 1,000 orders**, using AssemblyAI, Qwen3 4B (Ollama), Kokoro-82M, and FastAPI.
- Raised six-turn order accuracy from **0/12** to **10/10** by architecting a “model interprets, code decides” pipeline with schema-constrained intents and a deterministic resolver, keeping prompts flat at **~1.46k tokens** over **50 turns** via SQLite dialogue state.
- Cut LLM latency by **36.5%** (p50 **2.1s** → **1.3s**) and raised voice-session completion from **2/5** to **5/5**, eliminating dead air with **~50 ms** fillers and barge-in on AssemblyAI Universal-3.5 Pro streaming over WebSockets.

TECHNICAL SKILLS

AI & Machine Learning: LangGraph, LangChain, MCP, AssemblyAI, PyTorch, Scikit-learn, XGBoost

Backend & Databases: Python, FastAPI, PostgreSQL, Redis, ChromaDB, Qdrant

Frontend Development: TypeScript, React, Vite, HTML/CSS

Tools & Infrastructure: Docker, Git, Ollama, vLLM