

Long Quan Ton

HCMC, Vietnam | tonlongquanvn@gmail.com | [linkedin/in/benedictquanton](https://www.linkedin.com/in/benedictquanton) | [github/BenedictQuanTon](https://github.com/BenedictQuanTon) | [Portfolio](#)

EDUCATION

University of Technology Sydney & HCMUT

Ho Chi Minh City, Vietnam

Bachelor of Artificial Intelligence; GPA: 6.75/7.0 (Dean's List 2026); IELTS 7.0

Expected 2028

EXPERIENCE

FlyRank AI

Remote

AI Engineering Intern

Aug 2026 – Present

- Reduced new-workflow onboarding time by **60%** by designing a shared FastAPI and Redis layer that standardizes state, tool execution, and SQL access across LangGraph agents.
- Achieved **94%** valid text-to-SQL generation by building LangGraph multi-agent workflows on local LLMs (Ollama), with schema grounding and AST checks before execution.
- Cut mean time to resolve failed agent runs by **45%** by adding end-to-end request tracing and structured logging that ties every tool call to one trace ID.

Ho Chi Minh City University of Technology (HCMUT)

Ho Chi Minh City

Undergraduate AI Research Assistant – AITechLab (ML4U)

Aug 2026 – Present

- Researching **Automatic Speech Recognition (ASR)** in the Speech Recognition Team under Dr. Nguyen Duc Dung, analyzing model internals and benchmarking STT/TTS/LALMs systems for low-latency, real-time conversational voice agents, contributing to upcoming scientific publications and technical reports.
- Benchmarking STT, TTS, and large audio-language models (LALMs) for low-latency conversational voice agents, and evaluating Small Language Models (SLMs) for on-device speech-to-speech pipelines.

PROJECTS & COMPETITIONS

BKAi – Multi-Agent Admissions Counseling System

Personal Project

Full-Stack AI Developer

Jan 2026 – Apr 2026

- Built a LangGraph multi-agent RAG counselor with a deterministic verifier that answered **200/200** numeric admission questions exactly, by routing every score and quota to typed SQL tools over an 11-table fact DB and rejecting numbers absent from evidence.
- Lifted retrieval Hit@1 from 0.475 to **0.828** on 198 labeled queries by fusing Vietnamese dense and BM25 sparse vectors with RRF in Qdrant, plus cross-encoder reranking and parent-child chunking.
- Cut cold-answer p50 latency **15×** (27.3s → 1.79s) via a zero-LLM fast path, WebSocket streaming, and Gemini failover; shipped a Vietnamese voice agent at 1.26% CER with AssemblyAI STT and local Kokoro TTS.

Weatherise – Top 10 Finalist, Vietnam AI Open Hackathon

NVIDIA/Viettel/Sovico

Project Lead & AI Developer

Jun 2026

- Led a 4-person team to a **Top 10 finish** with a multi-agent weather decision system for Tourism, Construction & Agriculture, matching domain experts on **91%** of go/no-go calls (**120 scenarios**) at **4.8s median latency (1,247 runs)** on **8×** NVIDIA H200, using LangGraph, NeMo Agent Toolkit, and NVIDIA NIM.
- Cut forecast error by **12% MAE** vs. the best single source (**22%** vs. raw GFS, **180 days**) by fusing **7 live weather APIs** through a **10-tool** MCP Server with bias correction, outlier rejection, and an LLM arbiter, using FastAPI and Redis 7.
- Achieved **0 unsafe go-calls** across **212 red-team prompts** (94% blocked at input, the rest vetoed by a deterministic rule engine) and **87.2% context recovery** via a **4-collection** RAG layer (~**2,000 records**) behind NeMo Guardrails, using nv-embedqa-e5-v5, Qdrant, and PostgreSQL 16.

The Lantern – AssemblyAI Voice Agent Hackathon

AssemblyAI

Project Lead & AI Engineer

Sep 2026

- Led a 5-person team to build a real-time multilingual AI waiter (AssemblyAI, Qwen3 4B, Kokoro-82M, FastAPI) at under \$1 per 1,000 orders; a “model interprets, code decides” pipeline raised six-turn order accuracy from 0/12 to **10/10**.
- Cut LLM p50 latency by **36.5%** (2.1s → 1.3s) and raised voice-session completion from 2/5 to 5/5 test sessions using ~50ms fillers and barge-in over WebSocket streaming.

TECHNICAL SKILLS

AI & Machine Learning: LangGraph, LangChain, MCP, RAG, AssemblyAI, PyTorch, Scikit-learn, XGBoost

Backend & Data: Python, FastAPI, PostgreSQL, SQLite, Redis, Qdrant

Frontend: TypeScript, React, Vite

Tools & Infrastructure: Docker, Git, WebSockets, vLLM, Ollama

Certifications: Generative AI with LLMs (AWS/DeepLearning.AI), Machine Learning Specialization (Stanford/DeepLearning.AI), AWS Cloud Practitioner Essentials (AWS)