

Long Quan Ton

HCMC, Vietnam | tonlongquanvn@gmail.com | [linkedin/in/benedictquanton](https://www.linkedin.com/in/benedictquanton/) | [github/BenedictQuanTon](https://github.com/BenedictQuanTon) | [Portfolio](#)

EDUCATION

University of Technology Sydney & HCMUT

Bachelor of Artificial Intelligence; GPA: **6.75/7.0** (Dean's List 2026); IELTS 7.0

Ho Chi Minh City, Vietnam

Expected 2028

EXPERIENCE

FlyRank AI

AI Engineering Intern

Remote

Aug 2026 – Present

- Reduced new-workflow onboarding time by **60%** by designing a shared FastAPI and Redis layer that standardizes state, tool execution, and SQL access across LangGraph agents.
- Achieved **94%** valid text-to-SQL generation by building LangGraph multi-agent workflows on local LLMs (Ollama), with schema grounding and AST checks before execution.
- Cut mean time to resolve failed agent runs by **45%** by adding end-to-end request tracing and structured logging that ties every tool call to one trace ID.

Ho Chi Minh City University of Technology (HCMUT)

Undergraduate AI Research Assistant – AITechLab (ML4U)

Ho Chi Minh City

Aug 2026 – Present

- Co-authored **AAC-Bench** (under review, AAMAS 2027), a benchmark for proactive voice agents that use audio as context, in the Speech Recognition Team under Dr. Nguyen Duc Dung.
- Benchmarking STT, TTS, and large audio-language models (LALMs) for low-latency voice agents, and evaluating Small Language Models (SLMs) for on-device speech-to-speech pipelines.

PUBLICATIONS

AAC-BENCH: Benchmarking Audio as Context for Proactive Voice Agents

The Hieu Pham, Phuc Huynh, **Long Quan Ton**, Duc Dung Nguyen

AAMAS 2027

Under Review

PROJECTS & COMPETITIONS

BKAi – Multi-Agent Admissions Counseling System

Full-Stack AI Developer

Personal Project

Jan 2026 – Apr 2026

- Built a LangGraph multi-agent RAG counselor with a deterministic verifier that answered **200/200** numeric admission questions exactly, by routing every score and quota to typed SQL tools over an 11-table fact DB and rejecting numbers absent from evidence.
- Lifted retrieval Hit@1 from 0.475 to **0.828** on 198 labeled queries by fusing Vietnamese dense and BM25 sparse vectors with RRF in Qdrant, plus cross-encoder reranking and parent-child chunking.
- Cut cold-answer p50 latency **15×** (27.3s → 1.79s) via a zero-LLM fast path, WebSocket streaming, and Gemini failover; shipped a Vietnamese voice agent at 1.26% CER with AssemblyAI STT and local Kokoro TTS.

Weatherise – Top 10 Finalist, Vietnam AI Open Hackathon

Project Lead & AI Developer

NVIDIA/Viettel/Sovico

Jun 2026

- Led a 4-person team to a Top 10 finish with a multi-agent weather decision system (LangGraph, NeMo Agent Toolkit, NVIDIA NIM) that matched domain experts on **91%** of 120 go/no-go scenarios at 4.8s median latency.
- Cut forecast MAE by **12%** vs. the best single source by fusing 7 weather APIs through an MCP server with bias correction and an LLM arbiter; produced 0 unsafe go-calls across 212 red-team prompts via NeMo Guardrails and a rule engine.

The Lantern – AssemblyAI Voice Agent Hackathon

Project Lead & AI Engineer

AssemblyAI

Sep 2026

- Led a 5-person team to build a real-time multilingual AI waiter (AssemblyAI, Qwen3 4B, Kokoro-82M, FastAPI) at under \$1 per 1,000 orders; a “model interprets, code decides” pipeline raised six-turn order accuracy from 0/12 to **10/10**.
- Cut LLM p50 latency by **36.5%** (2.1s → 1.3s) and raised voice-session completion from 2/5 to 5/5 test sessions using ~50ms fillers and barge-in over WebSocket streaming.

TECHNICAL SKILLS

AI & Machine Learning: LangGraph, LangChain, MCP, RAG, AssemblyAI, PyTorch, Scikit-learn, XGBoost

Backend & Data: Python, FastAPI, PostgreSQL, SQLite, Redis, Qdrant

Frontend: TypeScript, React, Vite

Tools & Infrastructure: Docker, Git, WebSockets, vLLM, Ollama

Certifications: Generative AI with LLMs (AWS/DeepLearning.AI), Machine Learning Specialization (Stanford/DeepLearning.AI), AWS Cloud Practitioner Essentials (AWS)